



Securing Password Storage

Increasing Resistance to Brute Force Attacks

-jOHN (Steven)
Internal CTO

 @m1splacedsoul

<http://mobro.co/m1splacedstache>



v0.4

Chandu Ketkar
Technical Manager
 @cketkar

Scott Matsumoto
Principal Consultant
 @smatsumoto

Audience: OWASP crowd mainly; Developer secondly.



History /etc/password

etc/password

```
root:0:0:EC90xWpTKCo
hjackman:100:100:KMEzyulaQQ2
bgoldthwa:101:101:Po2gweIEPZ2
jsteven:102:500:EC90xWpTKCo
msoul:103:500:NTB4S.iQhwk
nminaj:104:500:a2N/98VTt2c
```

- Circa 1973
- 'one-way' password encryption
- chmod a+r /etc/passwd
- DES took 1 sec per password

2



...bringing us to 2012

```
00000fac2ec84586f9f5221a05c0e9acc3d2e670  
0000022c7caab3ac515777b611af73afc3d2ee50  
deb46f052152cfed79e3b96f51e52b82c3d2ee8e  
00000dc7cc04ea056cc8162a4cbd65aec3d2f0eb  
00000a2c4f4b579fc778e4910518a48ec3d2f111  
b3344eaec4585720ca23b338e58449e4c3d2f628  
674db9e37ace89b77401fa2bfe456144c3d2f708  
37b5b1edf4f84a85d79d04d75fd8f8a1c3d2fbde  
00000e56fae33ab04c81e727bf24bedbc3d2fc5a  
0000058918701830b2cca174758f7af4c3d30432  
000002e09ee4e5a8fcdae7e3082c9d8ec3d304a5  
d178cbe8d2a38a1575d3feed73d3f033c3d304d8  
00000273b52ee943ab763d2bb3d83f5dc3d30904
```

What do you see here?

How do we know what it is?

How could we figure this out?

In the news

LinkedIn

IEEE

Yahoo

...

SHA1('password') = 1e4c9b93f3f0682250b6cf8331b7ee68fd8

3



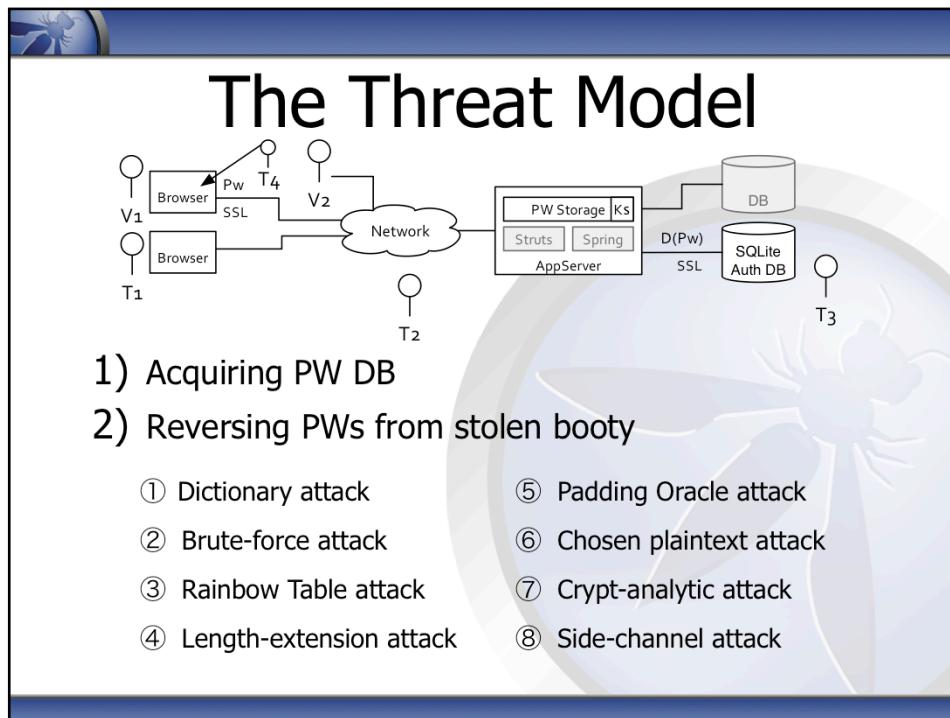
Golden Rules

- #1 – Don't be on the front page of InfoWeek
- #2 – Have a great story when you're on the front page of InfoWeek

Your passwords
WILL be
extracted from
your system

The Rules are based on quotes by the CEO of Merrill Lynch when Merrill brought their online banking system. Rule #1 was "Don't be on the front page of the Wall Street Journal". Rule #2 was "Don't be on the front page of the New York Times".

Here the second rule addresses the fact that we should assume that every vulnerability in the system/application cannot be plugged and that the password table will get into the wild.



This talk describes the control when the password table is leaked.

Introduce the problem in terms of the threat model – the Threat Model will be used later to help show the defense-in-depth

Describe Diagram of the Threat Model

- You can try to prevent an injection or something that extracts the passwords
- You have to assume that they will get out
- We're going to focus this talk on solutions for thwarting: T1 and T3 [T4 is outside of this discussion, but described below]
- See –jOHN's PW Storage Green field solution for a complete treatment on this diagram. T1 – T3 are excepted below
- The difference in terminology between this preso and jOHN's is that T3 is referred to as “The Internal Threat” rather than “LAN-based threat”

The following actors participate in this threat model:

[V1] - *Active System User* : Compromises one's self through use of the system
 Accesses the system normally through a browser
 May access the system through a compromised network (exposing them to



- Plaintext
- Encrypted
- Hashed (using SHA)
- Salt and Hash
- Adaptive Hashes
 - PBKDF
 - bcrypt
 - scrypt

Current Industry Practices

Engage the audience to see what they do. Gauge which people (companies) promote which solution.

The main point of this section is to look at Salt and Hash VS Adaptive Hash

The meta-point about this section is that all of the solutions in this section focus on a single control-point.

At the end of the section the point will be that Adaptive Hash

- Is more CPU intensive than a more conventional hash and will thus have a negative impact on scalability
- The notion of slowing down the attacker's ability to brute force is commendable, but it's debatable whether you can find a point that slows down the attacker enough while still maintaining reasonable CPU efficiency for (concurrent) logins.
- Thus, the hash-and-encrypt solution becomes a "bird in the hand" and the adaptive hash is a "two in the bush".



Hash Properties

`digest = hash(plaintext);`

- Uniqueness
- Determinism
- Collision resistance
- Non-reversibility
- Non-predictability
- Diffusion
- Lightning fast



First to understand the Properties of the Crypto function. The properties will restrict/influence your design.

Use a Better Hash?

 SHA-1

SHA-2

SHA-224/256

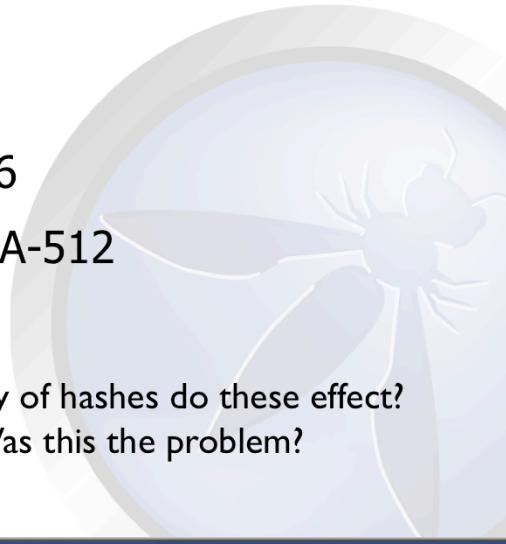
SHA-384/SHA-512

SHA-3

What property of hashes do these effect?

Collisions. – Was this the problem?

No.



If SHA1 is breakable, then just use a stronger (bigger) hash.

The different variants of SHA-2 vary the output size. The max message size is larger for SHA-384/SHA-512.

If you believe that there is a collision-based attack (two plaintext hashing to the same value), then you won't believe this

Can We Successfully Attack a Hash?

Depends on the threat-actor...

- Script-kiddie
- AppSec Professional
- Well-equipped Attacker
- Nation-state

Is the algorithm supported by a
tool?

Reprise or make the point about Threats Actor = Requirements.

For an unsalted hash (which is what an NTLM password is), we can use pre-computed rainbow tables.

Background on Rainbow tables (from Chandu). Note that this is just to get you started in understanding.

---- Begin From Chandu ---

Rainbow tables are generated as "optimized lookup tables" to reverse engineer one-way hash functions.

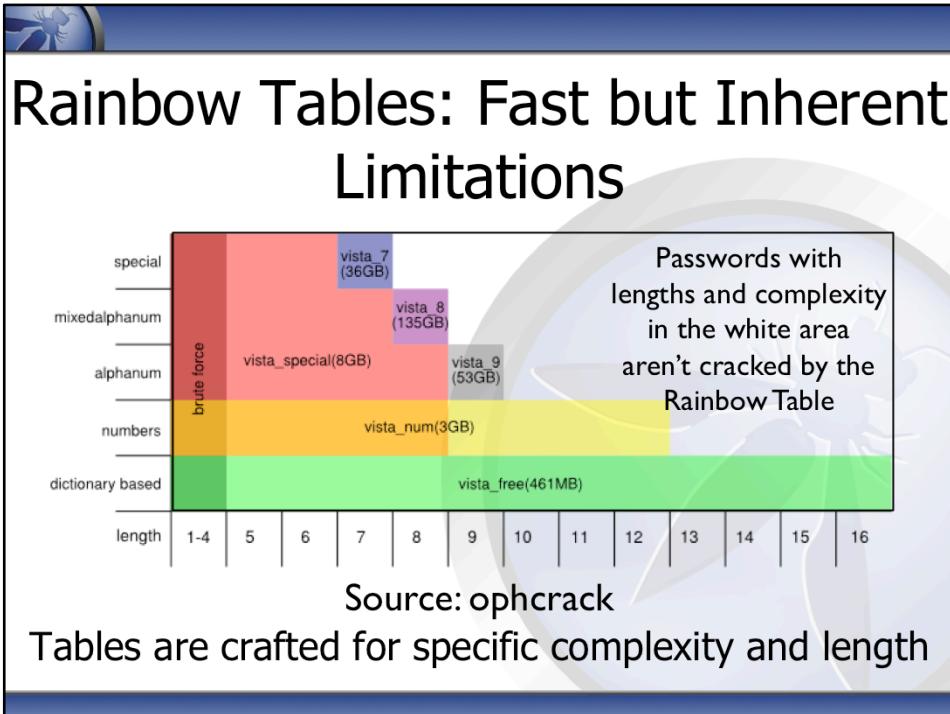
Key idea is this -

If I want to reverse engineer a 8-character password (alphanumeric – using upper and lower case + digits 0-9), then I would generate all possible 8-character strings and store these strings and their hashes in a table.

If I have a hash, then I can simply do a lookup in this table to find the plaintext password.

Rainbow tables are optimized

An 8-char password using lower/upper case letters (26+26) and digits (10), will require us to create (52^p8) a table with few trillion rows. You can compute 52^p8.



This graphic is from the ophcrack project on source forge. <http://ophcrack.sourceforge.net/tables.php>. What it shows is that RB tables have to be crafted with inherent limitations of size and valid-character space. Most of the tables are 99% accurate which means that not all passwords can be cracked with the table.

But, how many passwords do I need? All of them or just some of them?



Table Sizes

Search Space	Lookup Table (Brute Force)	Rainbow Table (NTLM hashes)
307,000 word dictionary	16 MB	461 MB
(a-z A-Z 0-9) ⁴	338 MB	8.0 GB
(a-z A-Z 0-9) ⁵	21 GB	8.0 GB
(a-z A-Z 0-9) ⁶	1.3 TB	8.0 GB
(a-z A-Z 0-9) ⁷	87 TB	8.0 GB
(a-z A-Z 0-9) ⁸	5,560 TB	134.6GB
(a-z A-Z 0-9) ⁹	357,000 TB	No table
(a-z A-Z 0-9) ¹⁰	22,900,149 TB	No table

Lookup table generates all possible combinations. The RB tables are about 99% accurate, but are limited in terms of the sizes and complexity of the passwords. They are, however, faster and most space efficient than a pure look up table.

“No Table” means that ophcrack doesn’t have a pre-calculated table – doesn’t mean that one can’t be generated

What Does the Salt Do?

```
salt || digest = hash(salt || plaintext);
```

De-duplicates digest texts

Adds entropy to input space*

- increases brute force time
- requires a unique table per user



<What is the difference between “brute force” and “rainbow table”? Seems like the solution for breaking the salted hash is always a table and the question is whether once can pre-compute the table and store it or whether a per-user-table must be generated on the fly.>



Can salted hashes be Attacked?

Depends on the threat-actor...

- Script-kiddie
- Some guy
- Well-equipped Attacker
- Nation-state

Attacking a table of salted hashes means building a Rainbow Table per user



For a Salted Hash we have to use dynamically generated tables: one per salt.

How “well-equipped” is “well-equipped”?

Per User Table Building

Brute Force Time for SHA-1 hashed,
mixed-case-a alphanumeric password

		8 Characters	9 Characters
Attacking a single hash (32 M/sec)	NVS 4200M GPU (Dell Laptop)	80 days	13 years
Attacking a single hash (85 M/sec)	\$169 Nvidia GTS 250	30 days	5 years
Attacking a single hash (2.3 B/sec)	\$325 ATI Radeon HD 5970	1 day	68 days

These numbers were generated using InsiderPro's PasswordPro password recovery program. Prices are from Amazon on 10/4/2012.

The point is that with cheap hardware, it's possible to get generating enough hashes per second to realistically brute force a password table protected with salted-SHA1. If you combine using a Rainbow Table tool which uses fewer hashes the numbers only get better.

Justin White generated the data on his laptop. Without getting into a treatise about number of instructions needed to execute the hash operation and intricacies about which operations are available on the processor... All of the performance data is anecdotal. There are some clips that I found below, but the point is that it's doable and with more money an attacker can construct a machine to using off-the-shelf hardware and tools from the Internet.

THe ATI Radeon HD is not at the top performer. It's an acceptable performing card. (http://www.videocardbenchmark.net/high_end_gpus.html)

---- Anecdotal discussions about how many hashes can be done per second ---
I have once made some experiences with SHA-1. A simple password hash with SHA-1 has the cost of processing a single "block" (SHA-1, like MD5, processes data by 64-



Algorithms designed specifically to remove the “lightning-fast” property of hashes

Thus: protecting passwords from Brute Force and Rainbow Table attacks

Adaptive Hashes increase the amount of time each hash takes through iteration

Adaptive Hashes

For an on-line attack, we can thwart T1 by increasing the time between tries. We can “Wait” for 1 second.

Adaptive Hashes are an off-line attack thwarting mechanism, but the problem is that there is no “wait”; there is only increased computation time.

“The added computational work makes [password cracking](#) much more difficult, and is known as [key stretching](#). “ – wikipedia (PBKDF2)

PW-Based Key Derivation (PBKDF)

```
salt || digest = PBKDF(hmac, salt, pw, c=);
```

Application Code:

```
salt = random.getBytes(8)
c = 10000000

key = pbkdf2(salt, pw, c, )
protected_pw = concat(salt, key)
```

Underlying implementation:

```
pbkdf2(salt, pw, c, b){
    r = computeNumOutputBlock(b)
    md[1] = SHA1-HMAC(p, s || 1)
    for (i=2; i <= c; i++)
        md[i] = SHA1-HMAC(p, md[i-1])

    for (j=0; j < b; j++)
        kp[j] = xor(md[1] || md[2] ... md[j])

    dk_b = concat(kp[1] || kp[2] ... kp[r])
    return dk_b
}
```

Well-supported & vetted

HMAC key is password

Attacker has all entropy

What is the right 'c'?

- NIST: 1000
- iOS4: 10000
- Modern CPU: 10000000

***SIMPLIFIED Code: see [IEEE RFC2898](#) for details
See [Java JCE Documentation](#) for details on Java API

- PBKDF says that it can be any Pseudo Random Function (PRF), but the implementation of PBKDF2 only supports HMAC-SHA1
- Signature Parameters:

PRF is a pseudorandom function of two parameters with output length *hLen* (e.g. a keyed [HMAC](#))

Password is the master password from which a derived key is generated

Salt is a [cryptographic salt](#)

c is the number of iterations desired

dkLen is the desired length of the derived key (length in bits)

DK is the generated derived key

- Invocation Actuals

- HMAC-SHA-1 – (only PRF implemented by default Java JCE)
- password – password
- salt – salt
- iOS4 uses 10,000
- 160 – SHA-1 generates 160 bit hashes

NOTE that PBKDF was designed to generate a key for password-based crypto; the implication is that it's not designed to be executed at scale. From RFC 2898:

“A key derivation function produces a derived key from a base key and other parameters. In a password-based key derivation function, the base key is a

The slide has a blue header bar with a small logo on the left. The main title is 'bcrypt' in a large, bold, black font. Below the title is a line of code: `c || salt || digest = bcrypt(salt, pw, c=);`. This line is partially obscured by a thick horizontal red bar.

Application Code:

```

salt = bcrypt.genSalt(12)
c = 10000000

c, salt, key = bcrypt(salt, pw, c)
protected_pw = concat(c, salt, key)

```

Underlying implementation:

```

bcrypt(salt, pw, c){
    d = "OrpheanBeholderScryDoubt"
    keyState = EksBlowfishSetup(c, salt, pw)

    for (int i=0, i < 64,i++){
        d = blowfish(keyState, d)
    }

    return c || salt || d
}

```

Not supported by JCE
 2^{cost} iterations slows hash operations
 Is 2^{12} enough these days?
 What effect does changing cost have on DB?
 Outputting 'c' helps
 Resists GPU parallelization, but not FPGA

FPGA – Field Programmable Gate Array – In short programmable hardware.
 $2^{**12} = 4096$

Parallelism is the key thing for the attacker.

BACKGROUND/REFERENCE:

1. Niels Provos and David Mazières, *The OpenBSD Project* <http://static.usenix.org/events/usenix99/provos/provos.pdf>

1. A bit unfathomable, but this pseudo-code was enlightening. cost is used in the setup of Blowfish. It's a power of 2

```

bcrypt(cost, salt, pwd)
state = EksBlowfishSetup(cost, salt, key)
ctext = "OrpheanBeholderScryDoubt"
repeat (64)
    ctext = EncryptECB(state, ctext)
    return Concatenate(cost, salt, ctext)
EksBlowfishSetup(cost, salt, key)
state = InitState()
state = ExpandKey(state, salt, key)

```



scrypt

salt || digest = scrypt(salt, pw, N, r, p, dkLen);

Application Code:

```
N = 16384
r = 8
P = 1
Key=scrypt(salt, pw, N, p, dkLen) {
protected_pw = concat(salt, key)
```

Underlying implementation:

```
scrypt(pw, salt, N, p, c) {
    for (i=0, i < pl, i++)
        b[i] = PBKDF2(pw, salt, 1, p*Mflen)
    for (i=0, i < pl, i++)
        b[i] = ROMmix(b[i], N)
    return PBKDF2(pw, b[1]||b[2]||...b[p-1], 1, dkLen)
}

MF(b, N) {
    x = b
    for (i=0, i < N-1, i++)
        v = /* Chain BlockMix(x) over N*/
    for (i=0, i < N-1, i++)
        j= /* Integrify(b) mod N */
        x = /* Chain BlockMix(x xor v[j]) */
    return x
}

BlockMix(r, b) { /* Chain Salsa20(b) over r */ }
```

Packages emerging, well-trodden than bcrypt
 Designed to defeat FPGA attacks
 Configurable

- N = CPU time/Memory footprint
- r = block size
- P = defense against parallelism

***DRAMATICALLY SIMPLIFIED Code:
 See [scrypt by C. Percival](#)
 See [scrypt kdf-01, Josefsson](#) for spec.

scrypt is the next step in the “use more resources” line of password protection algorithms. Not a lot to say here except that it’s even slower than bcrypt.

Q to audience: Does that make it better? (This is a setup for the next slide)

From Wikipedia:

The scrypt function is specifically designed to hinder such attempts by raising the resource demands of the algorithm. Specifically, the algorithm is designed to use a large amount of memory compared to other password-based KDFs, making the size and the cost of a hardware implementation much more expensive, and therefore limiting the amount of parallelizing an attacker can use (for a given amount of financial resources).

BACKGROUND/REFERENCES

Parameters (from the Standard)

Input:

- P Passphrase, an octet string.
- S Salt, an octet string.
- r Block size parameter.
- N CPU/Memory cost parameter, must be larger than 1, a power of 2 and



Adaptive Hash Properties

Motivations

Resists most Threats' attacks

- Concerted (nation-state) can succeed w/ HW & time

Simple implementation

Scale CPU-difficulty w/ parameter*

Limitations

1. Top priority is convincing SecArch

- $C=10,000,000$ == definition of insanity
- May have problems w/ heterogeneous arches

2. API parameters ($c=$) != devops

- Must have a scheme rotation plan

3. Attain asymmetric warfare

- Attacker cost vs. Defender cost

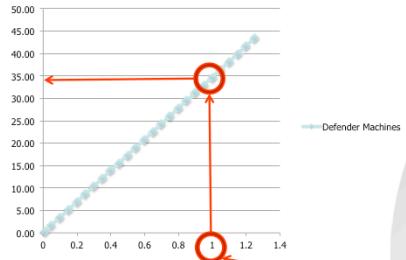
4. No password update w/o user

DEAL WITH MAINTAINANCE CPU, Architect, and other problems

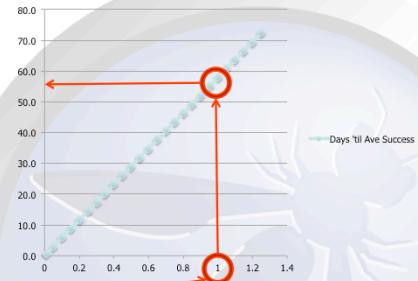
Defender VS Attacker	
Defender	Attacker
Goal: Log user in w/out > 1sec delay	Goal(s vary): Crack a single password, or <i>particular password</i> Create media event by cracking n passwords
Rate: 20M Users, 2M active / hr.	Rate: Scales w/ Capability
Burden: validation cost * users / (sec / hr.)	Burden: Bound by PW reset interval Population / 2 = average break = 10M
Hardware: 4-16 CPUs on App Server 2-64 servers	Hardware: Custom: 320+ GPUs / card, FPGA
Success Gauge : # of machines required for AuthN	Success Gauge: Days required to crack PW (ave)
Keep cost asymmetric: assure attacker cost greater than defender's	

Tradeoff Threshold

Machines Required to Conduct Login
(@ full load)



Days (average) Until Attacker Gets PW



- Is more than 8 AuthN machines reasonable?
- Is less than 2 months to average crack good enough?

Keep cost asymmetric: assure attacker cost greater than defender's



Adaptive Hashes At Best
Strengthen a Single
Control Point

We Can Do Better with
Defense In Depth

Requiring a Key
Gains Defense
In Depth



Hmac Properties

digest = hash(key, plaintext);

Motivations

Inherits hash properties

- This includes the lightning speed

Resists all Threats' attacks

- Brute force out of reach
 - $\geq 2^{256}$ for SHA-2
- Requires 2 kinds of attacks
 - AppServer: RMI Host keystore
 - DB: reporting, SQLi, backup

Limitations

1. Protecting key material challenges developers

- Must not allow key storage in DB!!!

2. Must enforce design to stop T3

- compartmentalization and
- privilege separation (app server & db)

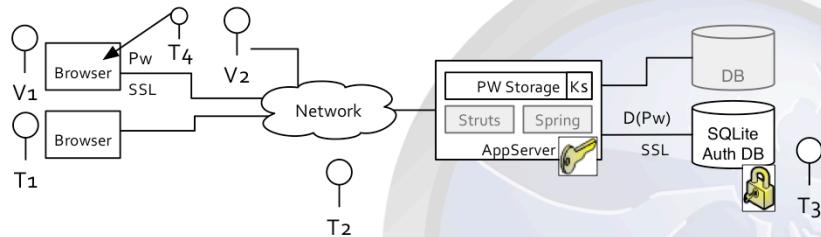
3. No password update w/o user

4. Stolen key & db allows brute force

- Rate \sim underlying hash function

COMPAT/FIPS Design

version||salt||digest = hmac(key, version||salt||password)



- Hmac = hmac-sha-256
- Version per scheme
- Salt per user
- Key per site
- Add a control requiring a key stored on the App Server
- Threats who exfiltrate password table also needs to get hmac key



Just Split the Digest?

No. They're not the same.

- Lacks key space (brute force expansion)
- Steal both pieces with the same technique
- Remember 000002e09ee4e5a8fcdae7e3082c9d8ec3d304a5 ?

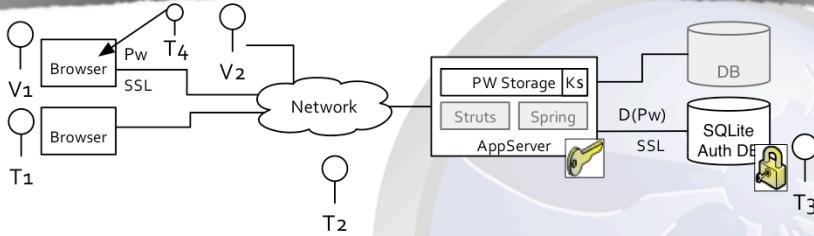
```
Permanence:code jsteven$ python split_hash_test.py -v 07606374520 -h ./hashes.txt
+ Found ['75AA8FF23C8846D1a79ae7f7452cfb272244b5ba3ce315401065d803'] verifying passwords
+ 1 total matching

Permanence:code jsteven$ python split_hash_test.py -h ./hashes_full.txt -v excal1ber -c 20
+ Found ['8FF8E2817E174C76b8597181a2ee028664aadff17a32980a5bad898c'] verifying passwords
+ 1 total matching
```

25

Reversible Design

```
version||cipher = ENC(wrapper keysite, <pw digest>)
<pw digest> = version||salt|| digest = ADAPT(version||saltuser||password)
```



- ENC = AES-256
- ADAPT = pbkdf2 | scrypt
- Version per scheme
- Salt per user
- Key per site

Reversible Properties

```
version||cipher = ENC(wrapper keysite, <pw digest>)
<pw digest> = version||salt|| digest = ADAPT(version||saltuser||password)
```

Motivations

- Inherits
 - “compat” solution benefits
 - Adaptive hashes’ slowness
- Requires 2 kinds of attacks
 - App Server & DB
 - Brute forcing DB out of reach ($>= 2^{256}$)
 - Stolen key can be rotated *w/o* user interaction
 - Stolen DB + key still requires reversing

Limitations

1. Protecting key material challenges developers
 1. Must not allow key storage in DB!!!
2. Must enforce design to stop T3
 1. compartmentalization and
 2. privilege separation (app server & db)
3. No password update w/o user
4. Stolen key & db allows brute force
 1. Rate \sim underlying adaptive hash



**MOST IMPORTANT
TOPIC**

Responding once attacked

Operations



Replacing legacy PW DB

1. Protect the user's account
 - Invalidate authN 'shortcuts' allowing login w/o 2nd factors or secret questions
 - Disallow changes to account (secret questions, OOB exchange, etc.)
2. Integrate new scheme
 - Hmac(), adaptive hash (scrypt), reversible, etc.
 - Include stored with digest
3. Wrap/replace legacy scheme: (incrementally when user logs in--#4)
 - `version||salt_new||protected = scheme_new(salt_old, digest_existing)` —or—
 - For reversible scheme: rotate key, version number
4. When user logs in:
 1. Validate credentials based on version (old, new); if old demand 2nd factor or secret answers
 2. Prompt user for PW change, apologize, & conduct OOB confirmation
 3. Convert stored PWs as users successfully log in

29



Thank You for Your
Time

Questions



Conclusions

- Without considering specific threats, the solutions misses key properties
- Understanding operations drives a whole set of hidden requirements
- Many solutions resist attack equivalently
- Adaptive hashes impose on defenders, affecting scale
- Leveraging design principles balances solution
 - Defense in depth
 - Separation of Privilege
 - Compartmentalization



TODO

- Revamp password cheat sheet
- Build/donate implementation
 - 1. Protection schemes
 - 2. Database storage
 - 3. Key store ← Vital to preventing dev err
 - 4. Password validation
 - 5. Attack response

32



Additional Material for
longer-format
presentations

Supporting Slides

Select Source Material

Trade material

[Password Storage Cheat Sheet](#)

[Cryptographic Storage Cheat Sheet](#)

[PKCS #5: RSA Password-Based Cryptography Standard](#)

[Guide to Cryptography](#)

Kevin Wall's [Signs of broken auth \(& related posts\)](#)

John Steven's [Securing password digests](#)

Graham-Cumming [1-way to fix your rubbish PW DB](#)

IETF [RFC2898](#)

Other work

[Spring Security, Resin](#)

[jascrypt](#)

Apache: [HTDigest](#), [HTTP Digest Specification](#), [Shiro](#)

Applicable Regulation, Audit, or Special Guidance

- COBIT DS 5.18 - Cryptographic key management
- Export Administration Regulations ("EAR") 15 C.F.R.
- [NIST SP-800-90A](#)

Future work:

- Recommendations for key derivation [NIST SP-800-132](#)
- Authenticated encryption of sensitive material: [NIST SP-800-38F \(Draft\)](#)

34



Threat Actors

Threat Actor	Attack Vector
[T1] External Hacker	AV0 - Observe client operations
	AV1 - Inject DB, bulk credentials lift
	AV2 - Brute force PW w/ AuthN API
	AV3 - AppSec attack (XSS, CSRF)
	AV4 - Register 2 users, compare
[T2] MiM	AV1 - Interposition, Proxy
	AV2 - Interposition, Proxy, SSL
	AV3 - Timing attacks
[T3] Internal/Admin	AV1 - Bulk credential export
	AV2 - [T1] style attack
	AV3 - Direct action w/ DB

These are the threat actors, but we only care about a subset of these Threats for the purpose of this talk.



Stored Passwords Requirements

Threat Actor	Attack Vector
[T1] External Hacker	AV0 - Observe client or network traffic
	AV1 - Inject DB, bulk credentials lift
	AV2 - Brute force w/ AuthN API
	AV3 - AppServer attack (XSS, CSRF)
[T2] MiM	AV4 - Register 2 users, compare
	AV1 - Interposition, Proxy
	AV2 - Interposition, Proxy, SSL
[T3] Internal/Admin	AV3 - Timing attacks
	AV1 - Bulk credential export
	AV2 - [T1] style attack
	AV3 - Direct action w/ DB
	Attack Vectors should be broken out by 1) acquisition of PW DB and 2) reversing the DB.

We're going to look at just the secure storage requirements part of the overall solution. Threat T2 and the grey-ed out attacks for T1 are in-scope for the overall application, but are not germane to the issue of storage. T1-AV1 and T3-AV1 are germane because it's through AV1 that the threat gets the password table.



COMPAT/FIPS Solution

```
<versionScheme>||<saltUser>||<digest> := HMAC(<keySite>, <mixed construct>
<mixed construct>      := <versionScheme>||<saltUser>||<pWUser>
• HMAC                  := hmac-sha256
• keySite                := PSMKeyTool(SHA256()):32B;
• saltUser               := SHA1PRNG():32B | FIPS186-2():32B;
• pWUser                 := <governed by password fitness>
```

Optional:

- <mixed construct> := <versionScheme>||<saltUser>||‘:’||<GUIDUser>||<pWUser>
- GUIDUser := *NOT* username or available to untrusted zones



hmac Solution Properties

Attack	Resistance
1.1 Resist chosen plain text attacks	Yes , Scheme complexity based on $(\text{salt}_{\text{user}} \& \text{pw}_{\text{user}}) + \text{key}_{\text{site}}$
1.2 Resist brute force attacks	Yes , $\text{Key}_{\text{site}} = 2^{256}$, $\text{salt}_{\text{user}} = 2^{256}$
1.3 Resist D.o.S. of entropy/randomness exhaustion	Yes , 32B on password generation or rotation
1.4 Prevent bulk exfiltration of credentials	Implementation detail: <various>
1.5 Prevent identical <protected>(pw) creation	Yes , provided by salt
1.6 Prevent <protected>(pw) w/ credentials	Yes , provided by Key_{site}
1.7 Prevent exfiltration of ancillary secrets	Implementation detail: store Key_{site} on application server
1.8 Prevent side-channel or timing attacks	N/A
1.9 Prevent extension, similar	Yes , hmac() construction (i_pad, o_pad)
1.10 Prevent multiple encryption problems	N/A (hmac() construction)
1.11 Prevent common key problems	N/A (hmac() construction)
1.12 Prevent key material leakage through primitives	Yes , hmac() construction (i_pad, o_pad)

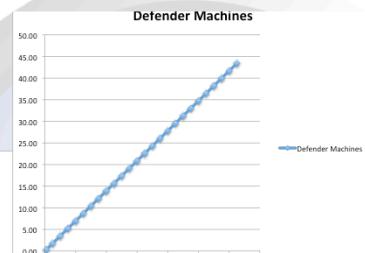
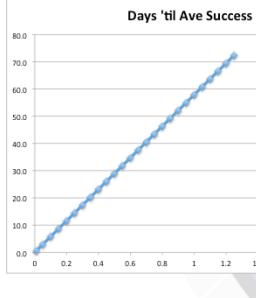
<describe the solution properties>

The problem is that it's difficult to rotate the key.

Attacker/Defender Worksheet

Attacker Speedup 2
 Average Pop. Yielding success 10000000
 Defender CPU 16
 Defender work (/ sec) 556

Seconds	Defender Machines	Days 'til Ave Success
0.01	0.35	0.6
0.05	1.74	2.9
0.1	3.47	5.8
0.15	5.21	8.7
0.2	6.94	11.6
0.25	8.68	14.5
0.3	10.42	17.4
0.35	12.15	20.3
0.4	13.89	23.1
0.45	15.63	26.0
0.5	17.36	28.9
0.55	19.10	31.8
0.6	20.83	34.7
0.65	22.57	37.6
0.7	24.31	40.5
0.75	26.04	43.4
0.8	27.78	46.3
0.85	29.51	49.2
0.9	31.25	52.1
0.95	32.99	55.0
1	34.72	57.9
1.05	36.46	60.8
1.1	38.19	63.7
1.15	39.93	66.6
1.2	41.67	69.4
1.25	43.40	72.3



39

(More) Just Split the Digest

Comparing 20B PBKDF2 chunks created no collisions

```
No spurious hit
Worst-case:
20B chunk = 50/50 split

• 2,150,710 uniquely
salted hashes
• 16 byte salt

• passwords
• mp3download
• REDROOSTER
• Dragon69
• 07606374520
• brazer1
• Bigwheel18
• Mastodon1
• Martha1a
• screaming36!
    Permanence: jsteven$ grep passwords ..[hashes.txt]
    Permanence: jsteven$ python split_hash_test.py -v passwords -h ..[hashes.txt]
    + Found [] matching passwords
    Permanence: jsteven$ python split_hash_test.py -h ..[hashes_full.txt] -v excaliber -c 20
    + Found 1 ['8FF8E2817E174C76b8597181a2ee028664adff17a32980a5bad898c'] matching passwords
    + Found 1 ['4F10C870B4E94F814fd07046b8d3bea650073e564c39596b8990d74b'] matching passwords
    + Found 1 ['EBD19B279CC64554f83f485706073fab5a112ea63143ec82a37e6d41'] matching passwords
    + Found 1 ['A4575F1E7D4C41Ec0ae49c5ce48ce4a9dbe28b9e87635e7289eb7eb'] matching passwords
    + Found 1 ['E1301662EC6349E5021c4cd8c158533aa9342ddee452f74f321ea0fa'] matching passwords
    + Found 1 ['72532DBFBF954FA1d9a068690ed1c3fc09459932be96bad5af4e1453'] matching passwords
    + Found 1 ['043EA3FE8434630d9d513284835c0891f0fbfcbeaf1f6bb6f76bc06'] matching passwords
    + Found 1 ['636BEF93F99449114785304641f419d450ce24ddfa03f4383e7593e6'] matching passwords
    + Found 1 ['A66772BEAF7A47361f6929611cc24b92b86cb84403c7773996ac49bc'] matching passwords
    + Found 1 ['8C8066C40C224A6700c50395afa1d3a87c9b76a1215193a29226e170'] matching passwords
    + Found 1 ['AD10E9DF1D23435163457052e8433cc60aaa853ee13143db9000456'] matching passwords
```

40